

Azure OpenAI Series (Virtual)

Saturday, 30 Mar, 2024

Unleashing the Power of Artificial Intelligence in the Cloud– Part 4



Aroh Shukla

Regional Microsoft Cloud Architect

Microsoft MVP Alumni, MCT



Aroh Shukla

MVP Alumni, MCT

Global Speaker

- Passionate to **learn**.
- Passionate to **share knowledge**.
- Passionate to work on **Microsoft Technologies**



Aroh.Shukla@gmail.com



/arohshukla



@aaroh_bits



AGENDA SLIDE

Part 1



01

AI Overview

what AI, responsible AI,
different types machine
learning models.



02

Computer Vision

Azure AI Vision,
included in Azure AI
Vision



03

NLP

conversational AI,
Azure AI services



04

Gen AI

Generative AI, Azure
OpenAI Service



05

Demos


All as discussed









What is Artificial Intelligence?

Software that **imitates human capabilities**

- **Predicting outcomes** and **recognizing patterns** based on historic data.
- Recognizing **abnormal events** and **making decisions**.
- Interpreting **visual input**.
- **Understanding language** and engaging in conversations.
- **Extracting information** from sources to gain knowledge.









Common AI workloads

	Machine Learning	Predictive models based on data and statistics – the foundation for AI.
	Computer Vision	Capabilities within AI to interpret the world visually through cameras, video, and images.
	Natural Language Processing	Capabilities within AI for a computer to interpret written or spoken language and respond appropriately. (e.g. ChatGPT, Bing Copilot)
	Document Intelligence	Capabilities within AI that deal with managing, processing , and using high volumes of data found in forms and documents.
	Knowledge Mining	Capabilities within AI to extract information from large volumes of often unstructured data to create a searchable knowledge store.
	Generative AI	Capabilities within AI that create original content in a variety of formats including natural language, image, code, and more .



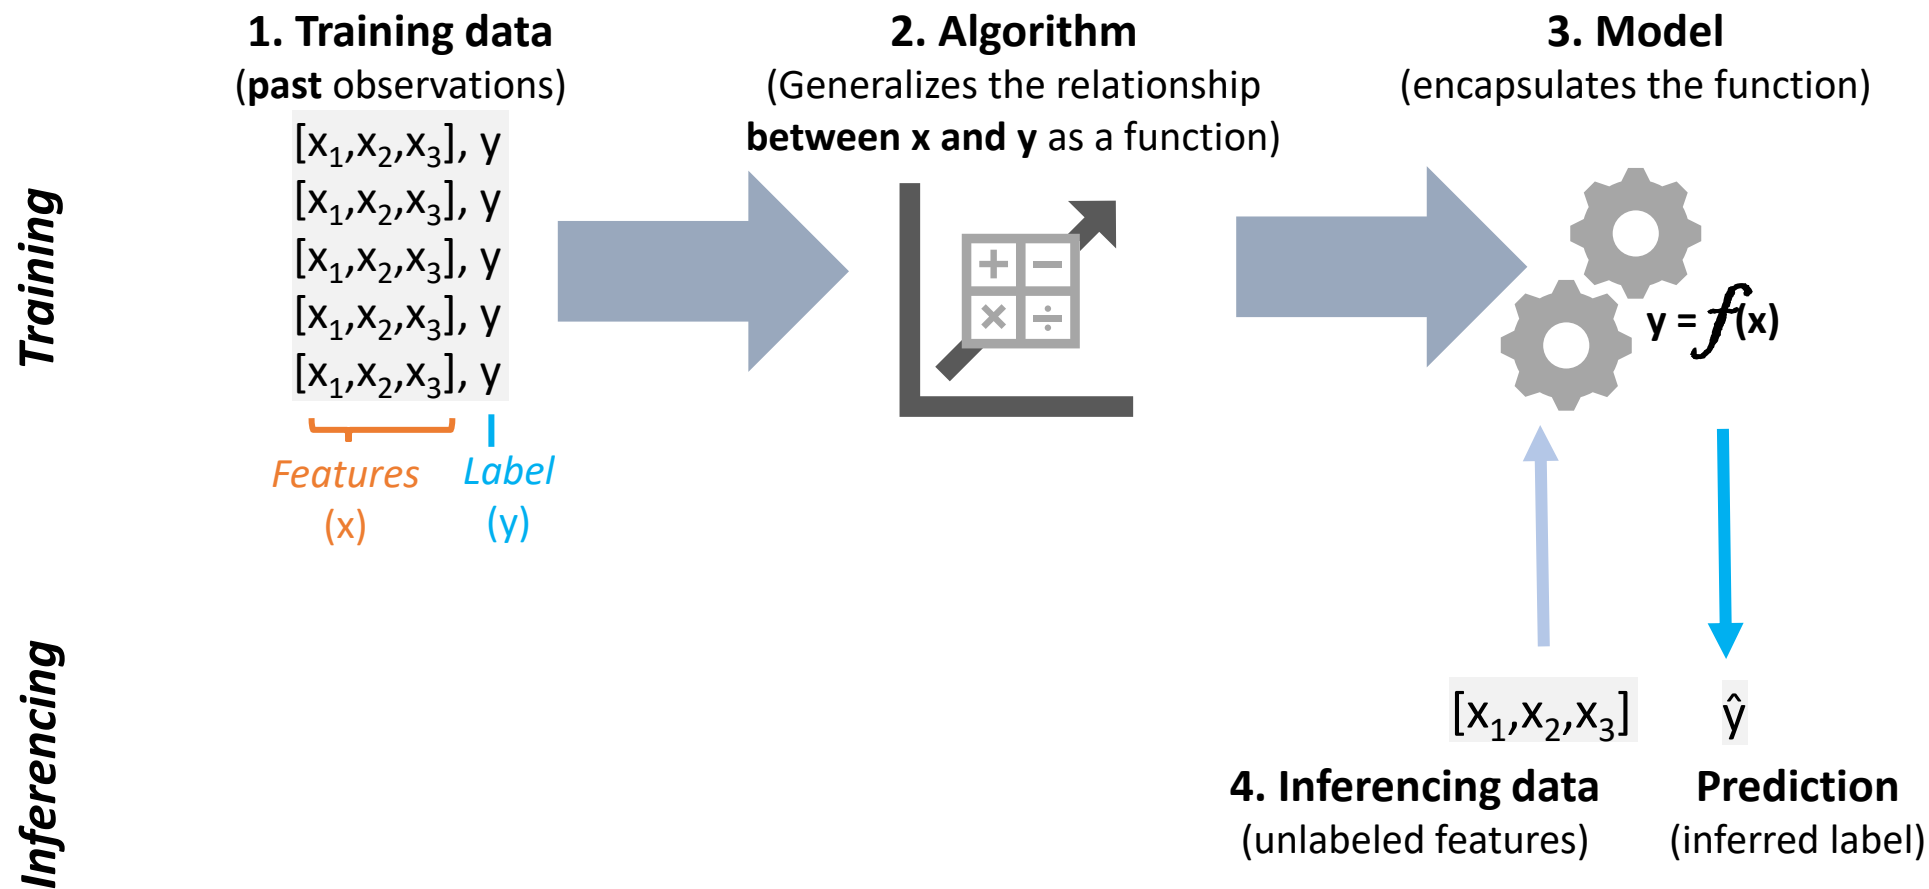
Principles of responsible AI

		Challenge or Risk	Example
	Fairness	Bias can affect results.	A loan-approval model discriminates by gender due to bias in the data with which it was trained.
	Reliability & safety	Errors may cause harm.	An autonomous vehicle experiences a system failure and causes a collision.
	Privacy & security	Private data could be exposed.	A medical diagnostic bot is trained using sensitive patient data , which is stored insecurely.
	Inclusiveness	Solutions may not work for everyone.	A predictive app provides no audio output for visually impaired users .
	Transparency	Users must trust a complex system.	An AI-based financial tool makes investment recommendations – what are they based on?
	Accountability	Who's liable for AI-driven decisions?	An innocent person is convicted of a crime based on evidence from facial recognition – who's responsible?



What is machine learning?

Creating predictive models by finding relationships in data



Types of machine learning

Machine Learning

Supervised machine learning


Training data includes known labels

Unsupervised machine learning

Training data is unlabeled

Regression

Label is a numeric value




Predict the **number of ice creams sold** based on day, season, and weather

Classification

Label is a categorization (or *class*)

Binary classification


Label *is* or *is not* a class



Predict whether a patient is at-risk **for diabetes based** on clinical data

Multiclass classification


Label is one of multiple classes



Predict the species of a **penguin based on its measurements**

Clustering

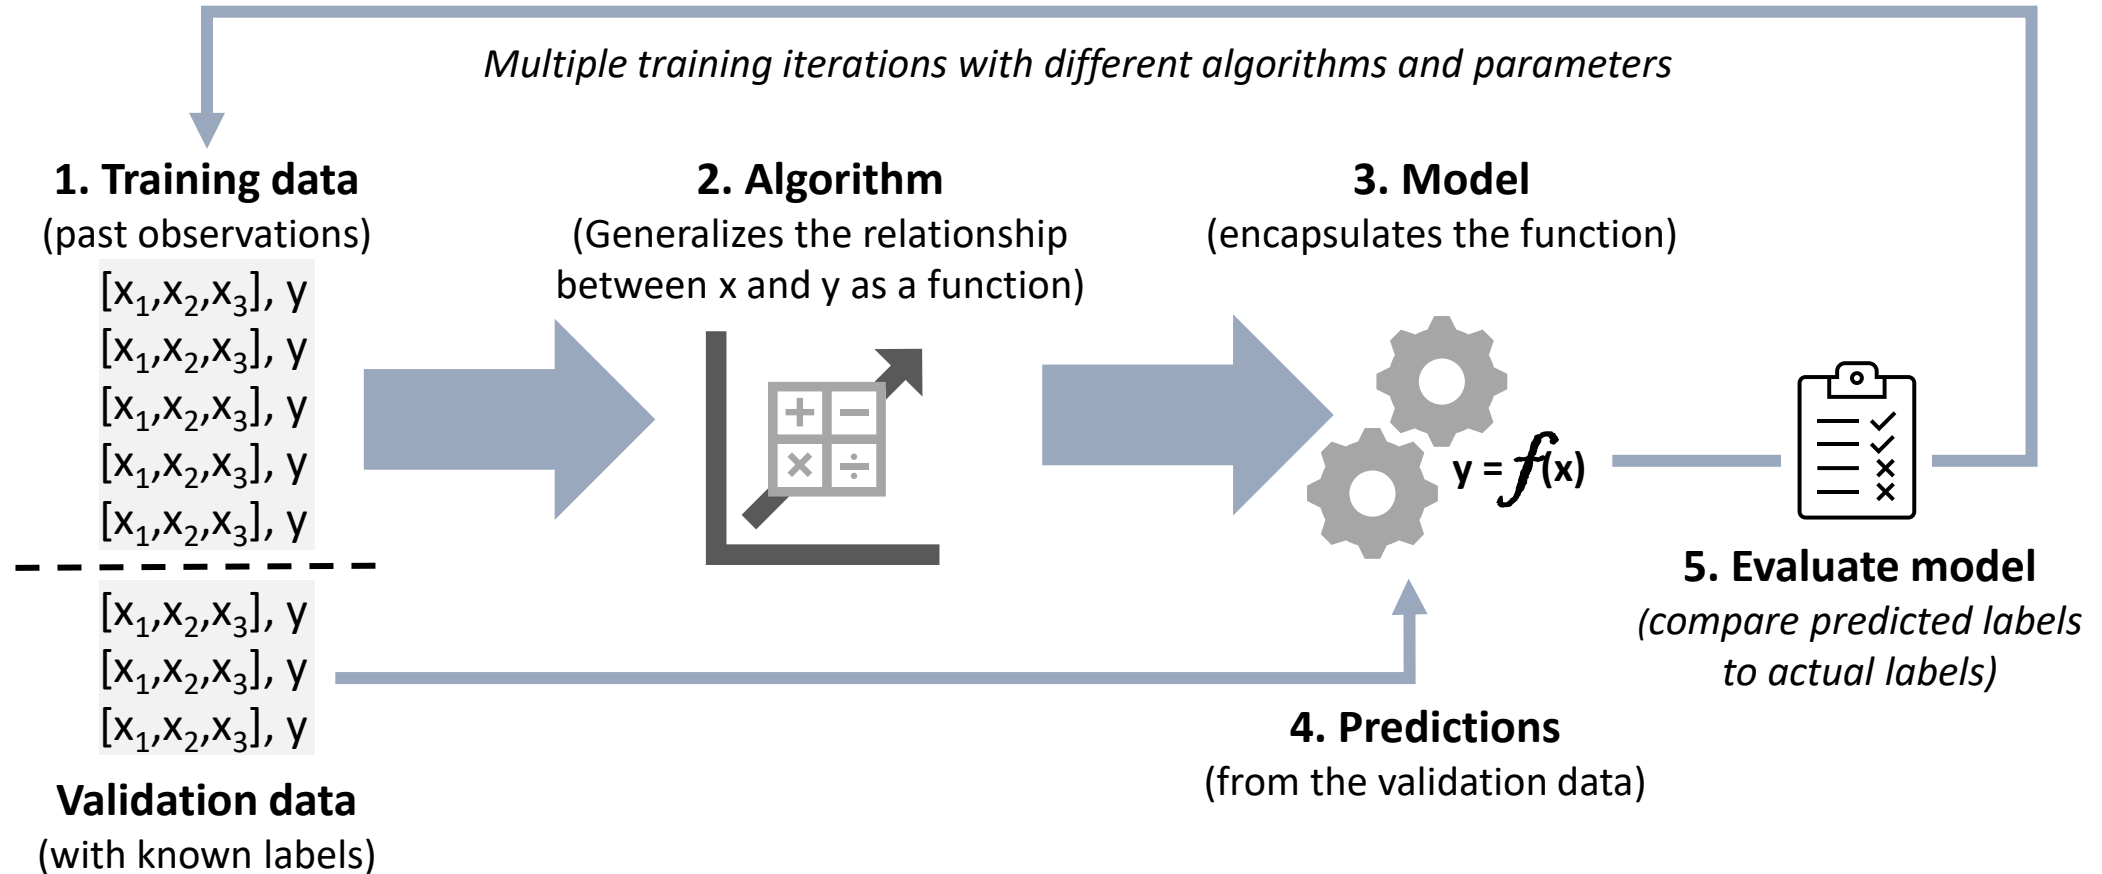
Similar items are grouped together



Separate plants into groups based **on common characteristics**



Model training and evaluation



Predicative AI vs Generative AI

	Predicative AI	Gen AI
Function	Predicts the most probable output given input data.	Creates new data similar to the input data distribution.
Training	Typically trained on labeled data for classification tasks.	Often trained on unlabeled data for generating new content.
Output Variety	Limited to predicting predefined classes or values.	Capable of producing diverse outputs, sometimes novel.




Deep learning (Subset of ML)

Human neural network



- **Neurons fire** in response to electrochemical stimuli
- **When fired**, the signal is passed to connected neurons

Artificial neural network

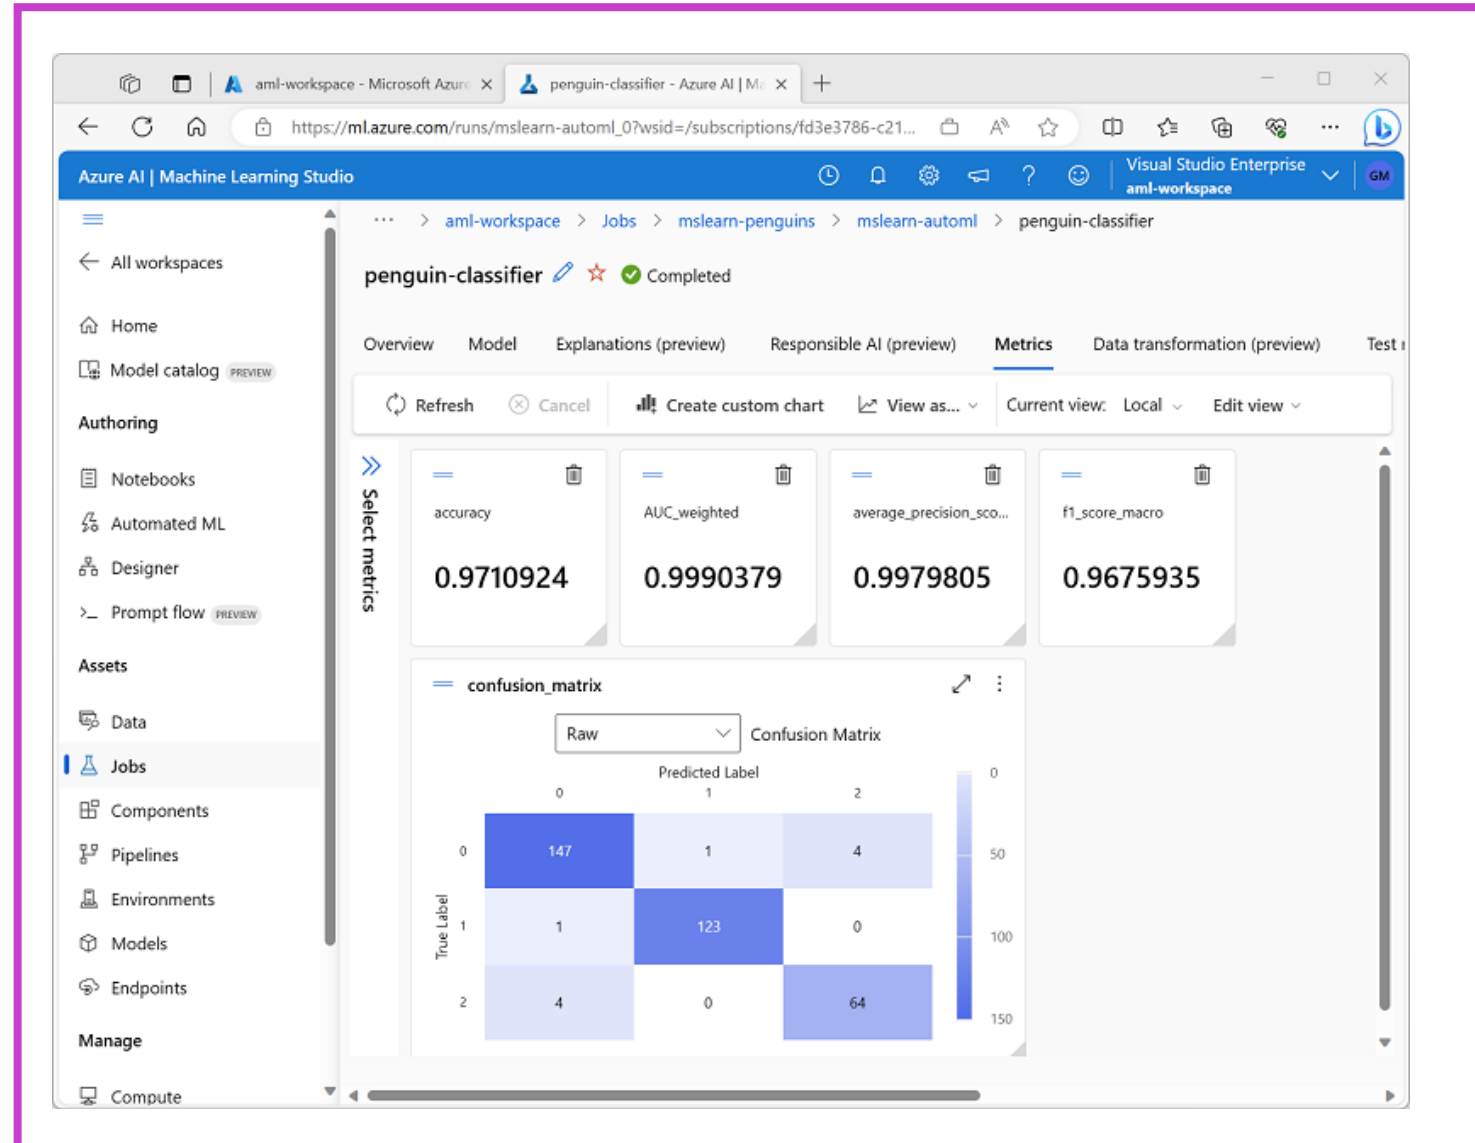


- Each neuron is a function that operates on an **input value (x)** and a **weight (w)**
- The function is wrapped in an **activation function** that determines whether to pass the output on



What is Azure Machine Learning (Azure ML)?

- **Azure Machine Learning** is a **cloud-based platform** for machine learning.
- **Azure Machine Learning Studio** is a user interface for accessing Azure Machine Learning capabilities.
- **Machine learning models** trained with Azure Machine Learning can be published as services.



Demo: Explore Automated Machine Learning in Azure Machine Learning Studio



In this demo, you will see how machine learning features can be used to train a machine learning model to make predictions.

1. Follow along on the exercise page at: <https://aka.ms/ai900-auto-ml>

Images and image processing

An image is an array of pixel values

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	255	255	255	0	0
0	0	255	255	255	0	0
0	0	255	255	255	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0


Filters are applied to change images

-1	-1	-1	0	0	0	0
-1	8	-1	0	0	0	0
-1	-1	255	255	255	0	0
0	0	255	0	255	0	0
0	0	255	255	255	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



Convolutional Neural Networks

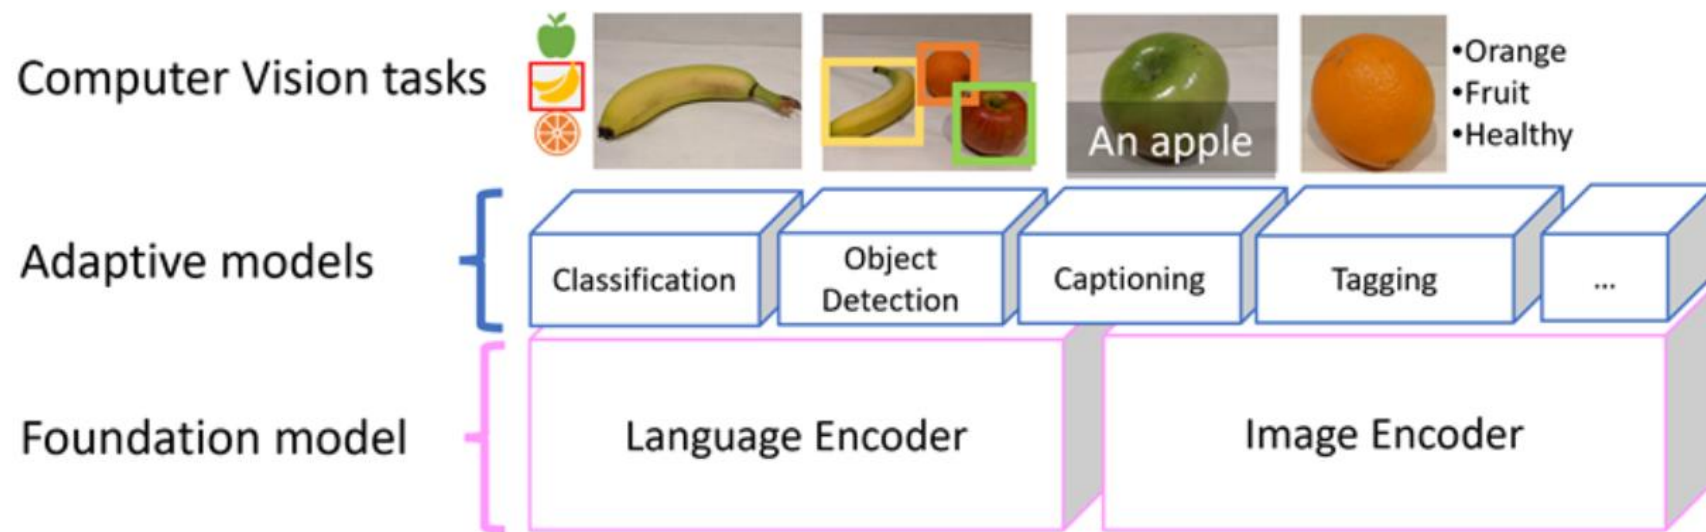
1. Labeled mages are used to train the model
2. Filter layers extract *feature maps* from each image
3. The feature maps are flattened
4. The feature values are fed into a fully connected neural network
5. The output layer produces a probability value for each possible class label



- During training, the filter kernels start with random weights. These weights are iteratively adjusted to improve the accuracy of the predictions based on the known labels.
- The trained model uses learned weights to extract features from new images and predict their class.



Multi-modal models



- A newer approach to modeling involves combining language and vision models that encode image and text data
- The model encapsulates semantic relationships between features extracted from the images and text extracted from related captions.
- A multi-modal model can be used as a *foundation* model for more specialized *adaptive* models.



Computer vision services in Azure



Vision

- Image Analysis:
 - Image tagging, captions, model customization, and more.
- Optical Character Recognition (OCR)
- Spatial analysis



Face

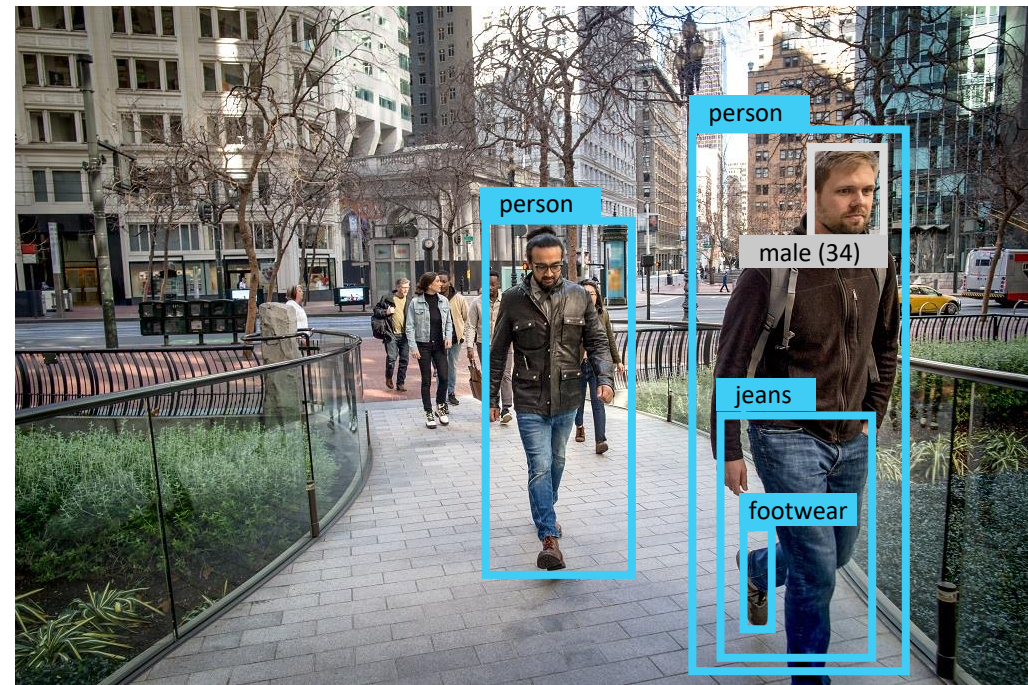
- Face detection
- Face recognition



Image analysis 4.0 with the *AI Vision* Service

Capabilities include:

- Model customization
- Read text from images
- Detect people in images
- Generate image captions
- Detect objects
- Tag visual features
- Smart crop



Caption: A group of people walking on a sidewalk

Tags: Building, jeans, street, outdoor, jacket, city, person



Detecting faces with the *Face* Service


Anyone can use the Face service to detect:

- Blur
- Exposure
- Glasses
- Head pose
- Noise
- Occlusion

Only Managed Microsoft customers can access facial recognition capabilities:

- Similarity matching
- Identity verification

*To support Microsoft's Responsible AI Principles, Facial Recognition is under a Limited Access policy.



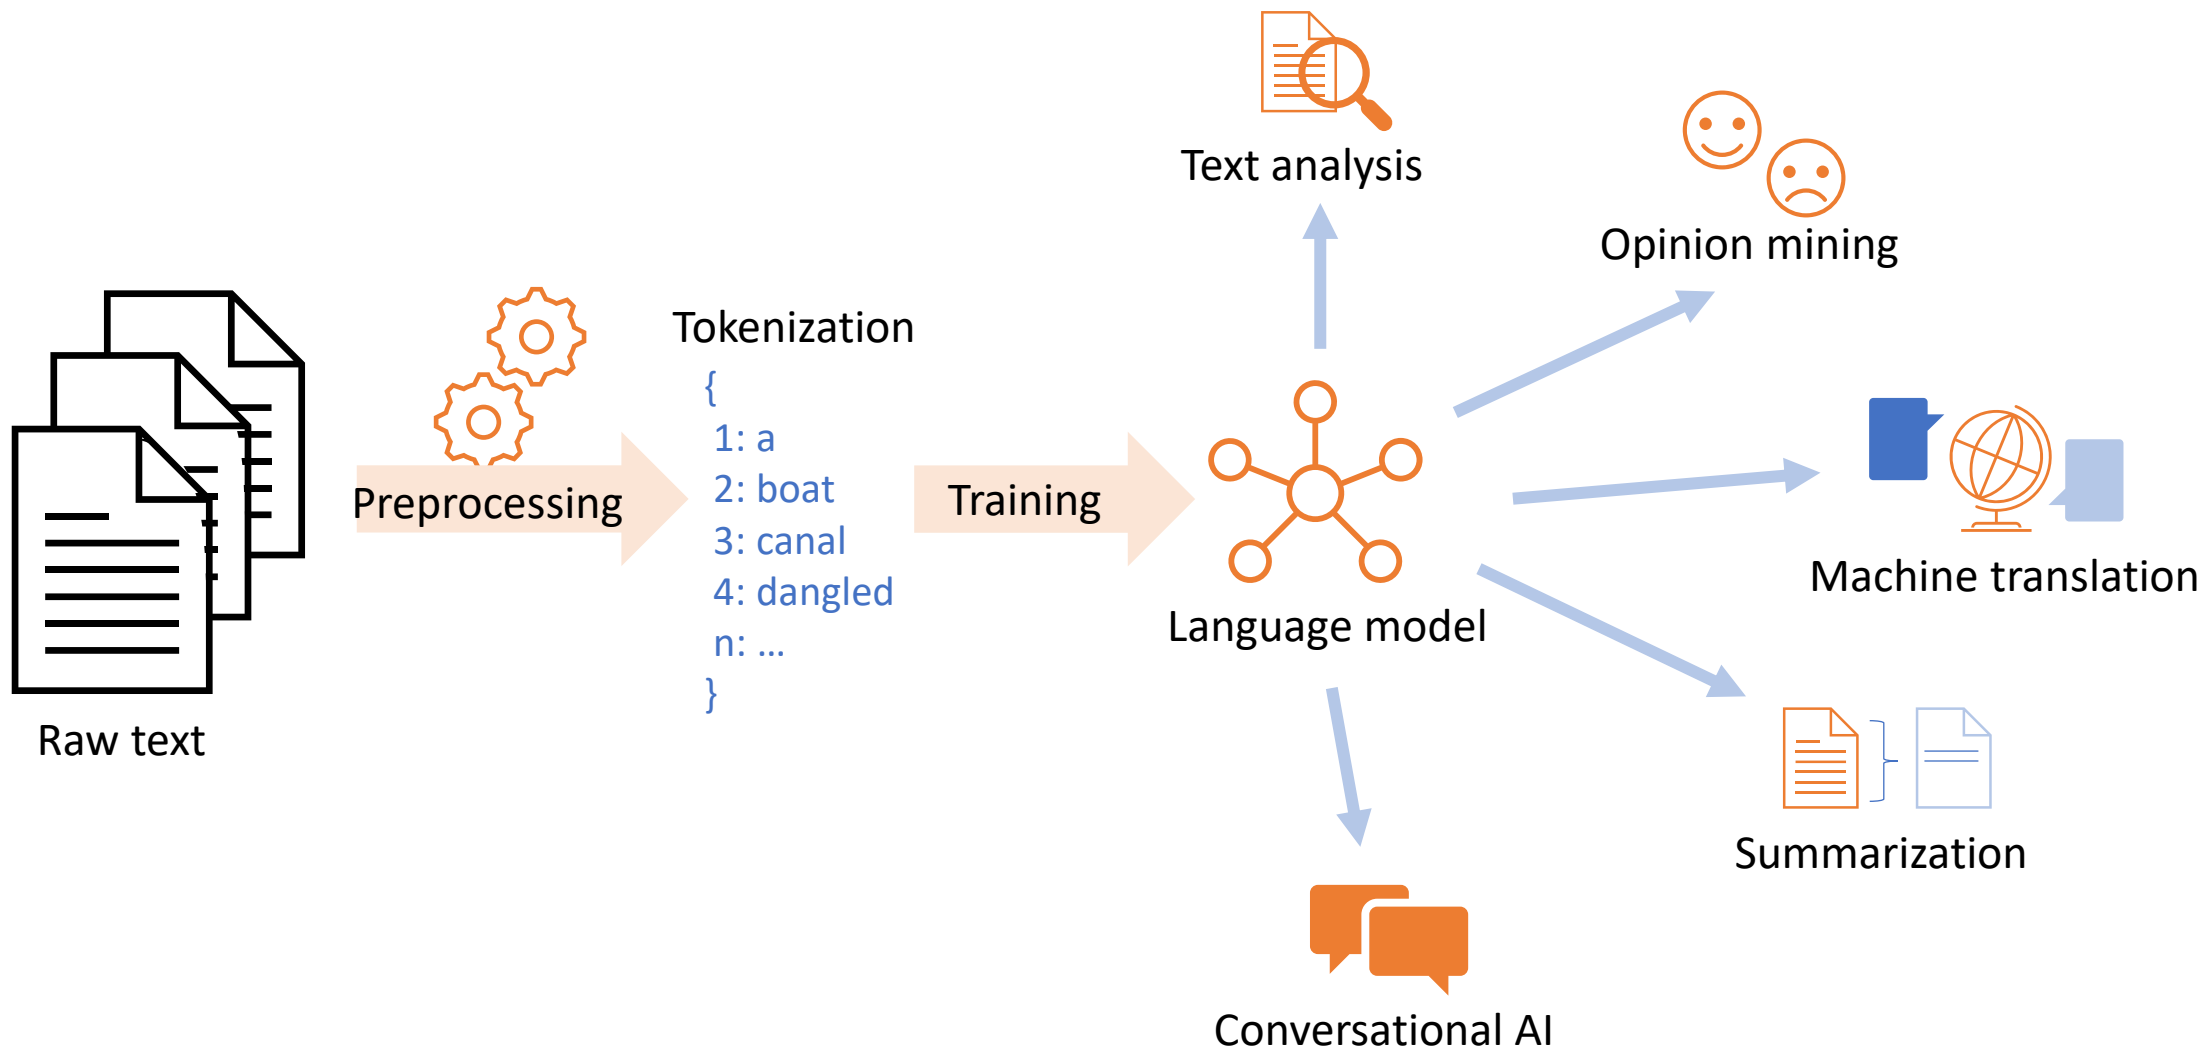
Demo: Detect faces in Vision Studio



In this demo, you will take a look at the face detection capabilities of the **Azure AI Face** service.

1. Follow along on the exercise page at: <https://aka.ms/ai900-face>

What is natural language processing?



Natural language processing and conversational AI in Azure



Language

- **Language detection**
- Key phrase extraction
- Named entity detection
- **Sentiment analysis and opinion mining**
- **Personal information detection**
- **Summarization**
- **Question answering**
- **Conversational language understanding**
- ...



Speech

- **Text to speech**
- **Speech to text**
- **Speech translation**
- Speaker identification
- Language identification
- ...



Translator

- **Text translation**
- **Document translation**
- Custom translation
- ...



Analyzing text

I had a wonderful vacation in France.

- **Predominant Language:** English
- **Sentiment:** .88 (positive)
- **Key Phrases:** “wonderful vacation”
- **Entities:** France

The screenshot displays the Azure AI Language Studio web interface. At the top, a blue header bar contains the text 'Azure AI | Language Studio'. Below this, the page title 'Language Studio' is followed by a 'Welcome to Language Studio' message. A section titled 'Recent custom projects you've worked on' includes a '+ Create new' button and a table with columns for 'Name' and 'Type'. Below the table is a 'View all projects' link. A navigation bar at the bottom of the main content area features four tabs: 'Featured', 'Extract information', 'Classify text' (which is selected and underlined), and 'Understand questions and conversational language'. The 'Classify text' tab displays three large, light-gray cards. The first card, 'Analyze sentiment and mine opinions', features an icon of a screen with a smiley face and a sad face, and describes detecting positive, negative, and neutral sentiment. The second card, 'Detect language', features an icon of a screen with language codes like {C}, {A}, and {J}, and describes evaluating text to detect a wide range of languages and dialects. The third card, 'Custom text classification', features an icon of a screen with a gear and a bar chart, and describes training a classification model. Each card has a 'Try it out' or 'Open Custom text classification' link at the bottom.




Question answering

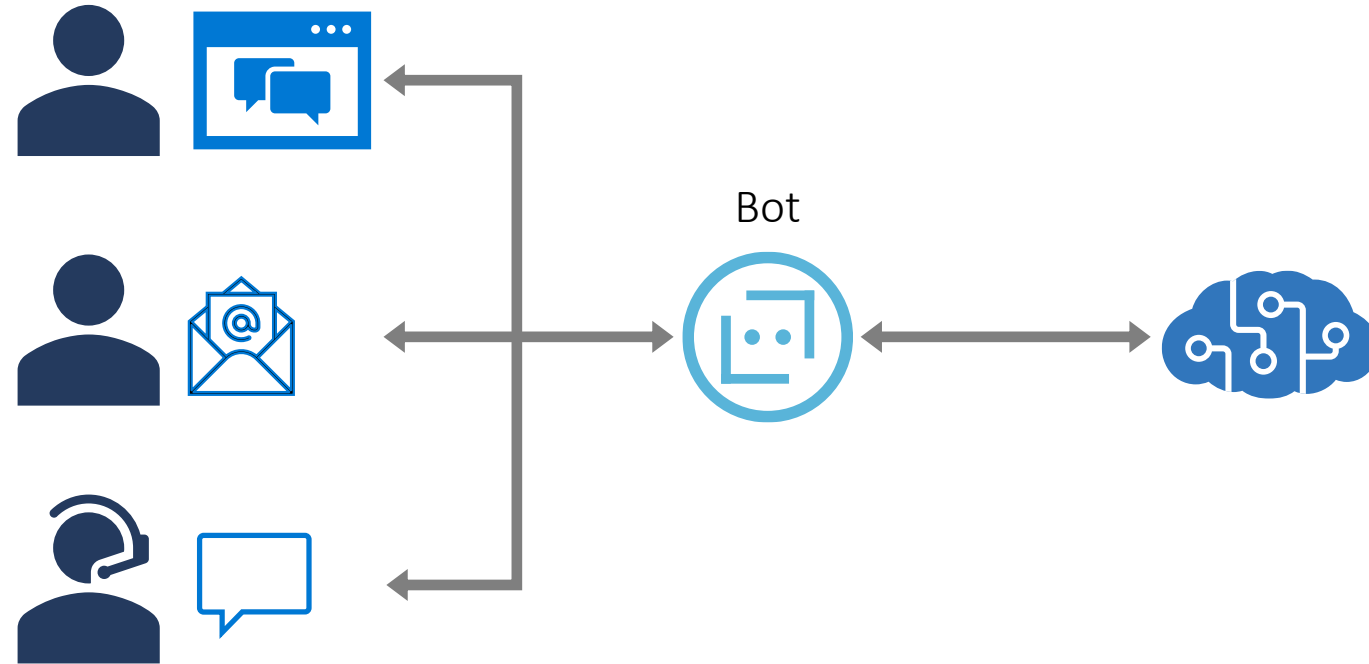
Define a *knowledge* base of **question and answer** pairs:

- By entering questions and answers
- From an **existing FAQ document**
- By using built-in *chit-chat*

Consume **the knowledge base** from client apps, **including bots.**



Azure bot service



- Cloud-based platform for developing and managing bots
- Integration with **AI** Language and other services
- Connectivity through multiple channels



Demo: Use Form Recognizer



In this demo, you'll see **Azure AI Language's** question answering capabilities in action.

1. Follow along on the exercise page at: <https://aka.ms/ai900-question-answering>

What is generative AI?

AI: imitates human behavior by using machine learning to interact with the environment and execute tasks without explicit directions on what to output.

Generative AI: creates original content, such as generative AI that has been built into chat applications. Generative AI applications take in natural language input, and return appropriate responses in a variety of formats:



Natural language generation



Image generation



Code generation



Large language models

Generative AI applications are powered by *large language models* (LLMs), which are a specialized type of machine learning model that you can use to perform natural language processing (NLP) tasks, including:

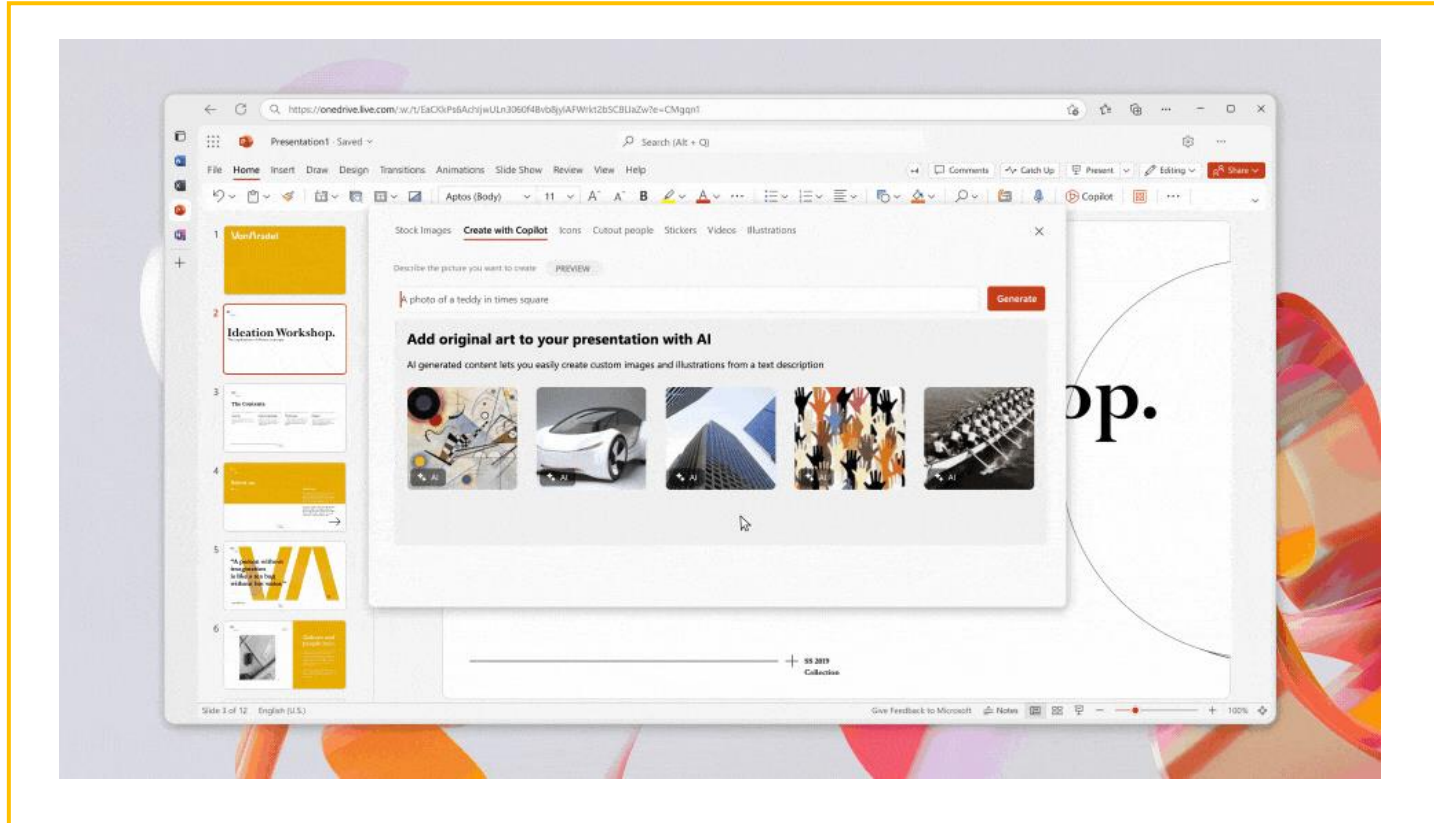
- Determining *sentiment* or otherwise classifying natural language text.
- Summarizing text.
- Comparing multiple text sources for semantic similarity.
- Generating new natural language.



Copilots

Copilots are often **integrated into other applications** and provide a way for users to get help with common tasks from a generative AI model.

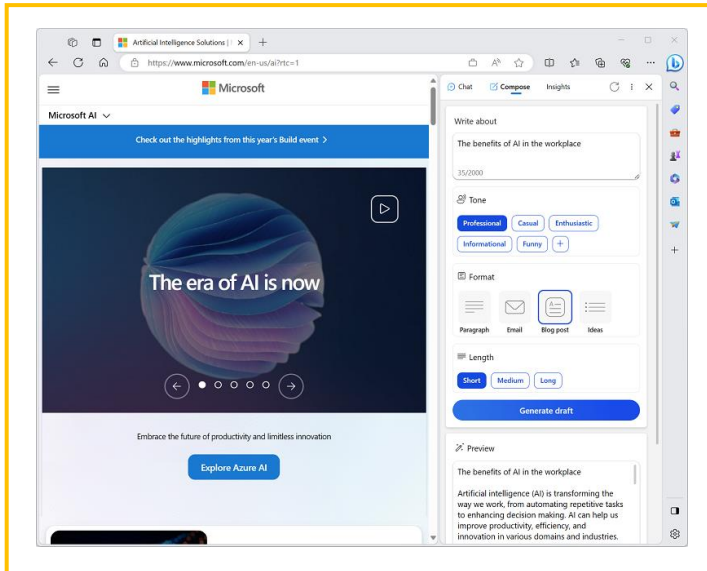
- **Developers can build copilots that submit prompts** to large language models and generate content for use in applications.
- **Business users can use copilots** to boost their productivity and creativity with AI-generated content.



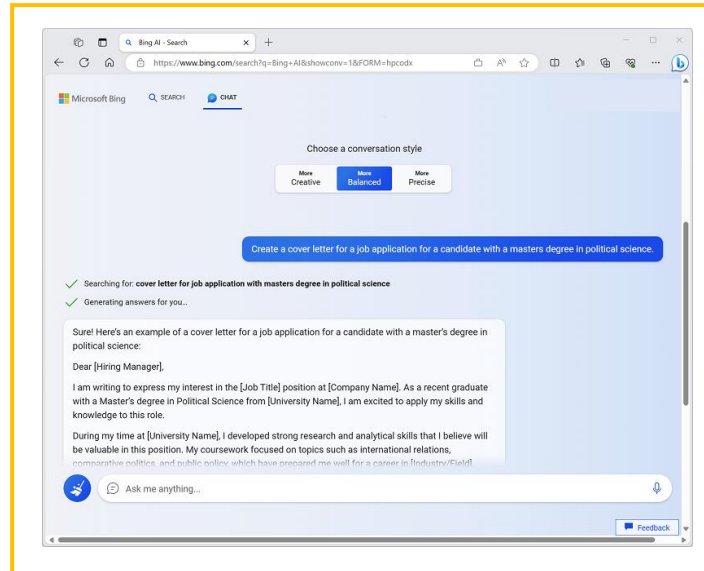
What are copilots?

Other examples of copilots

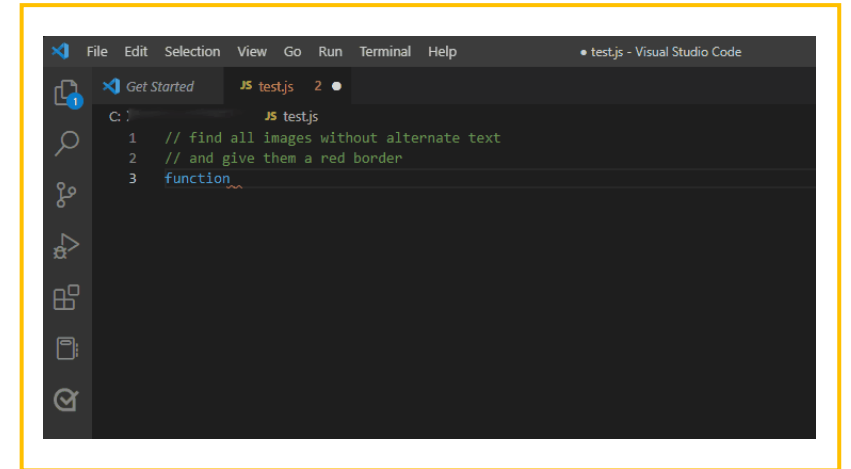
Microsoft Edge browser copilot



Microsoft Bing copilot



GitHub copilot



And many more!



Improve generative AI responses with prompt engineering

The term *prompt engineering* describes the process of prompt improvement.

Both developers who design applications and consumers who use applications can improve the quality of responses from generative AI by using direct language, system messages, examples, and/or grounding data.

	Description	Example
Direct language	You can get the most useful completions by being explicit about the kind of response you want.	“Create a list of 10 things to do in Edinburgh during August”.
System messages	Describe how the chat should act.	"You're a helpful assistant that responds in a cheerful, friendly manner“.
Providing examples	LLMs generally support <i>zero-shot learning</i> in which responses can be generated without prior examples. However, you can also provide a few example responses, known as <i>few-shot learning</i> .	“Visit the castle in the morning before the crowds arrive”.
Grounding data	You can include <i>grounding</i> data to provide context.	Including email text with the prompt “Summarize my email”.



What is Azure OpenAI?

Azure OpenAI service is Microsoft's cloud solution for deploying, customizing, and hosting large language models.

Azure OpenAI service consists of:

- Pre-trained generative AI models.
- Customization capabilities.
- Built-in tools to detect and mitigate harmful use cases so users can implement AI responsibly.
- Enterprise-grade security with role-based access control (RBAC) and private networks.

You can use several methods to develop Azure OpenAI solutions: Azure AI Studio, REST API, supported SDKs, and Azure CLI.



What models does Azure OpenAI support?

Azure OpenAI supports many LLMs:**Description**

GPT-4

A set of models that improve on GPT-3.5 and can understand as well as generate natural language and code.

GPT-3.5

A set of models that improve on GPT-3 and can understand as well as generate natural language and code.

Embeddings

A set of models that can convert text into numerical vector form to facilitate text similarity.

DALL-E (preview)

A series of models in preview that can generate original images from natural language.



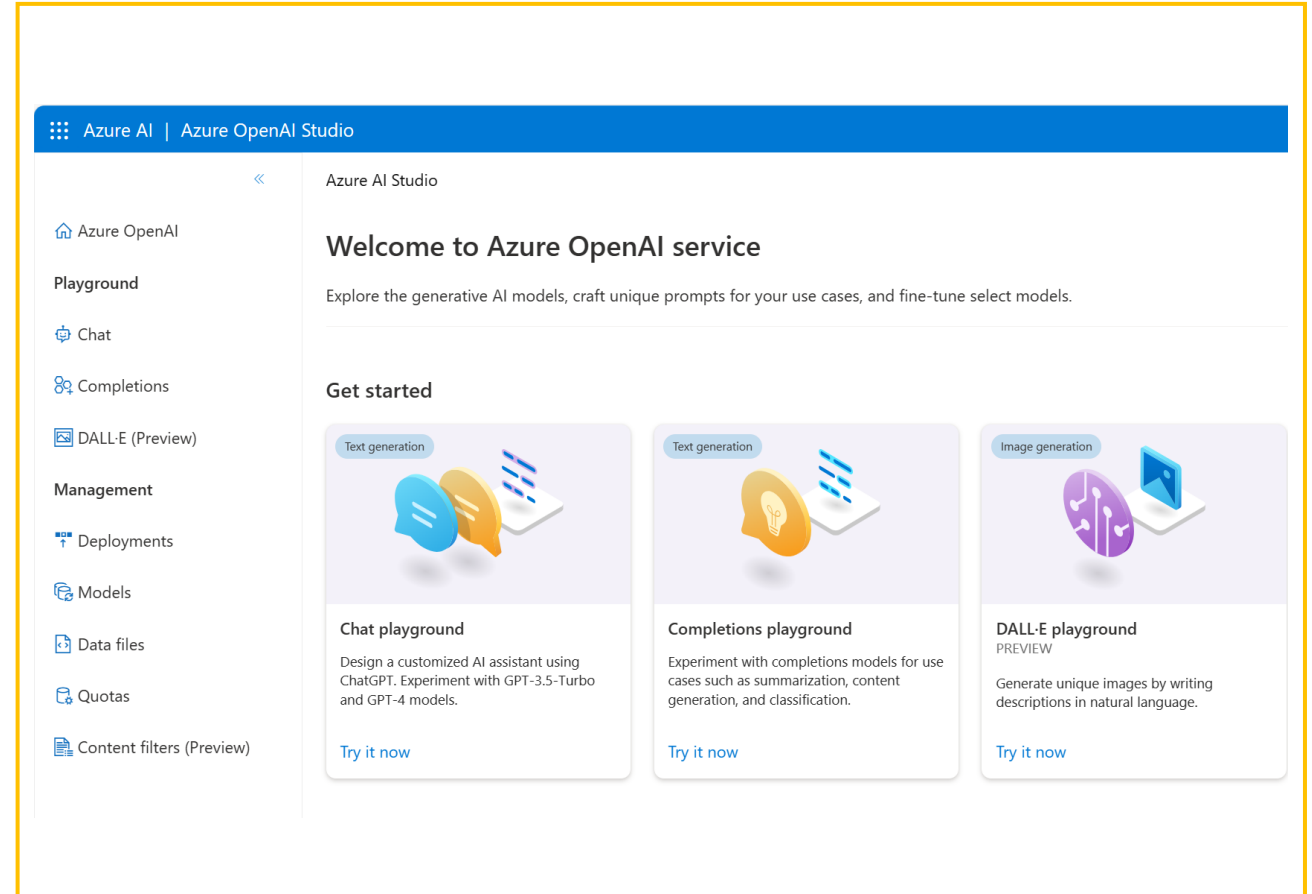
How to use Azure OpenAI

Azure OpenAI Studio:

- Build and deploy AI models for software applications
- Powered by generative AI models optimized for various tasks
- Azure OpenAI models include: **GPT-4**, **GPT-3.5**, **Embeddings** and **DALL-E models**

Playgrounds:

- Experiment with Azure OpenAI models without coding
- Use the assistant setup to instruct the model about how it should behave



Azure OpenAI's natural language capabilities

Generative pre-trained transformer (GPT) models are excellent at both understanding and creating natural language.

For example, given a prompt where the user types in text asking for a cooking recipe:

```
Write recipe instructions for a fruit pie  
based on these ingredients:
```

```
Strawberries  
Blueberries  
Flour  
Eggs  
Milk
```

Here's an example of the Azure OpenAI generated response:

```
Instructions:
```

1. Preheat your oven to 350 degrees Fahrenheit.
2. In a medium bowl, mix together 1 cup of flour and 1 egg until it forms a dough.
3. Grease a 9-inch pie dish with butter or non-stick cooking spray.
4. Roll out the dough and place it in the pie dish.
5. In a separate bowl, mix together 1 cup of strawberries, 1 cup of blueberries, and 1 cup of milk.
6. Pour the fruit mixture into the pie dish.
7. Bake



Azure OpenAI's code generation capabilities

GPT models translate natural language or code snippets into code.

Code generation goes beyond just writing code from natural language prompts. Given the following code, it can generate unit tests as shown on the right:

Python

```
# Python 3
def mult_numbers(a, b):
    return a * b

# Unit test
def
```



Python

```
# Python 3
def mult_numbers(a, b):
    return a * b

# Unit test
def test_mult_numbers():
    assert mult_numbers(3, 4) == 12
    assert mult_numbers(0, 10) == 0
    assert mult_numbers(4, 0) == 0

# Unit test
def test_mult_numbers_negative():
    assert mult_numbers(-1, 10) == -10
    assert mult_numbers(10, -1) == -10
```




Azure OpenAI's image generation capabilities

Generative AI models can edit and create images. The model that works with images is called DALL-E, which supports image creation, image editing and image variations creation.

- **Image generation:** With DALL-E, you can even request an image in a particular style. Styles can be used for edits and variations as well.
- **Editing an image:** DALL-E can edit the image as requested by changing its style, adding or removing items, or generating new content to add.
- **Image variations:** Image variations can be created by providing an image and specifying how many variations of the image you would like.

Prompt: Create four variations of an image of an elephant with a hamburger.



Demo: Explore Azure OpenAI Service



In this demo,
you'll see **Azure
OpenAI's**
capabilities in
action.

1. Follow along on the exercise page at:
<https://aka.ms/ai900-azure-openai>